

# 一、实验目的

理解网络中网络故障出现的必然性  
理解网络验证工具VeriFlow的原理  
掌握VeriFlow的检测网络故障的方法  
提高阅读工程代码、修改代码的能力

# 二、实验环境

实验在VMware运行的Arch Linux虚拟机上进行

组件	版本/配置
控制器	Ryu
虚拟化平台	Mininet
SDN网络验证工具	Veriflow

# 三、热身内容

## 观察现象

### 1. 启动拓扑

```
sudo ./simple.py
```

### 2. 启动控制程序

```
TOPO=simple.txt CONFIG=simple.config.json uv run ryu-manager  
ryu.app.ofctl_rest as_switch.py
```

### 3. 在拓扑中ucla2 ping purdue建立连接

```
mininet> ucla2 ping purdue
```

发现ping不通。查看控制程序，发生NameResoulutionError:Failed to resolve 'localhost' ([Errno -2] No address found)  
而启动控制程序时提示

```
wsgi starting up on http://0.0.0.0:8080
```

报错信息提示

```
File "/home/sdn/sdn-lab4/utlis/flowmod.py", line 35, in send_flow_mod
```

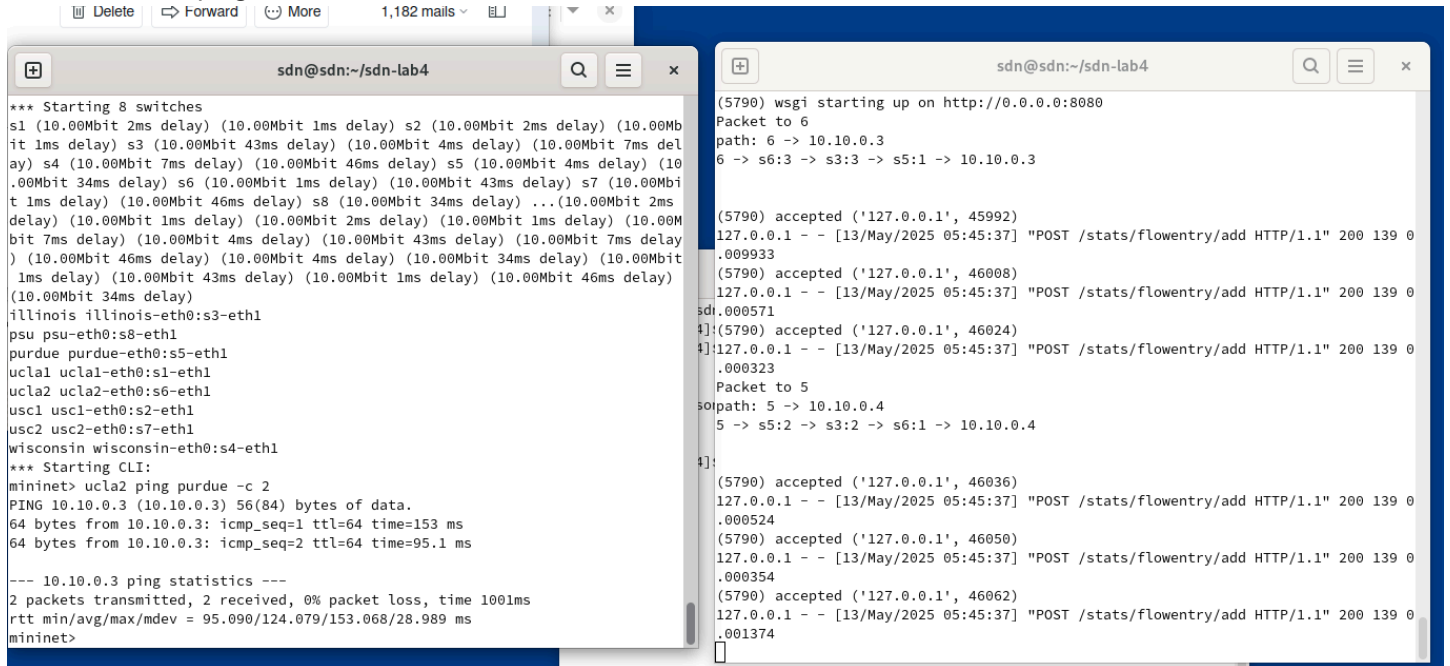
打开[flowmod.py](#),找到错误位置。

```
url = "http://localhost:8080/stats/flowentry/add"
```

修改为

```
url = "http://0.0.0.0:8080/stats/flowentry/add"
```

再次尝试,发现ping通,建立连接。

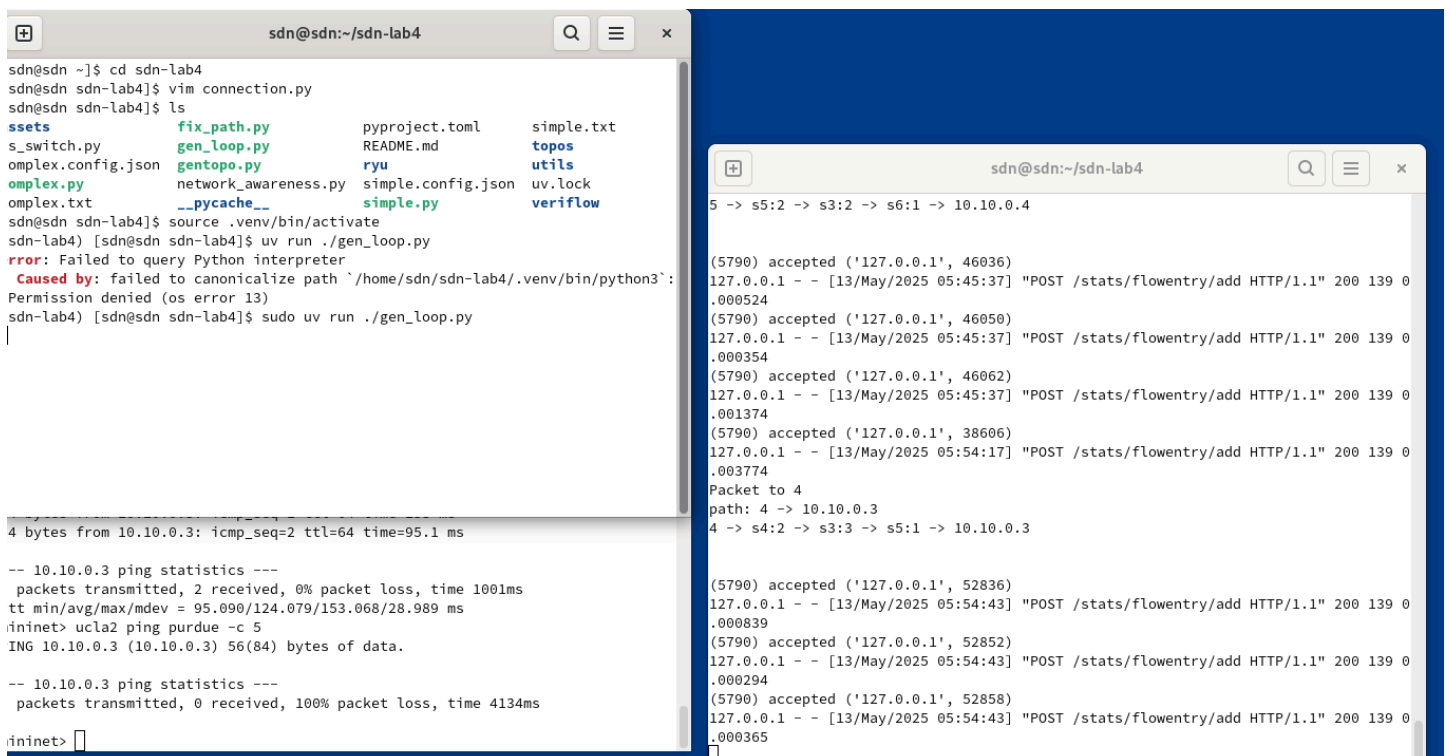


The image shows two terminal windows. The left window, titled 'sdn@sdn:~/sdn-lab4', displays the output of a Mininet CLI command 'mininet> ucla2 ping purdue -c 2'. It shows a successful ping with 64 bytes of data, an ICMP sequence of 1, and a TTL of 64. The right window, titled 'sdn@sdn:~/sdn-lab4', shows a log of HTTP POST requests to the '/stats/flowentry/add' endpoint, indicating that the flow entry was successfully added for the path 6 -> 10.10.0.3.

#### 4. 接着下发从illinois途经wisconsin到达ucla2的路径之后, 尝试ucla2 ping purdue失败

```
uv run ./gen_loop.py #下发从illinois途经wisconsin到达ucla2的路径
```

再次尝试ucla2 ping purdue,发现失败。



The image shows two terminal windows. The left window, titled 'sdn@sdn:~/sdn-lab4', displays the output of a 'ls' command, showing a directory listing of files and folders. The right window, titled 'sdn@sdn:~/sdn-lab4', shows a log of HTTP POST requests to the '/stats/flowentry/add' endpoint, indicating that the flow entry was successfully added for the path 5 -> s5:2 -> s3:2 -> s6:1 -> 10.10.0.4.

## 5. 查看路径上某一个交换机，如illinois的流表，发现匹配某一条流表的数据包数目异常增加

```
sudo ovs-ofctl dump-flows s3
```

```
(sdn-lab4) [sdn@sdn sdn-lab4]$ sudo ovs-ofctl dump-flows s3
cookie=0x0, duration=174.321s, table=0, n_packets=5, n_bytes=490, priority=10,ip,nw_dst=10.10.0.0/16 actions=output:"s3-eth4"
cookie=0x0, duration=693.768s, table=0, n_packets=2, n_bytes=196, priority=5,ip,nw_src=10.10.0.3,nw_dst=10.10.0.4 actions=output:"s3-eth2"
cookie=0x0, duration=148.384s, table=0, n_packets=0, n_bytes=0, priority=5,ip,nw_src=10.10.0.4,nw_dst=10.10.0.3 actions=output:"s3-eth3"
cookie=0x0, duration=700.312s, table=0, n_packets=28, n_bytes=1960, priority=0 actions=CONTROLLER:65509
(sdn-lab4) [sdn@sdn sdn-lab4]$
```

## 使用VeriFlow验证环路问题

### 1. 编译VeriFlow.

```
make -C veriflow -j$(nproc)
```

### 2. 在自定义端口开启远程控制器，运行AS控制器程序

```
TOPO=simple.txt CONFIG=simple.config.json uv run ryu-manager
ryu.app.ofctl_rest as_switch.py --ofp-tcp-listen-port 1024
```

### 3. 运行VeriFlow的proxy模式

将simple.txt移入到veriflow文件夹中。

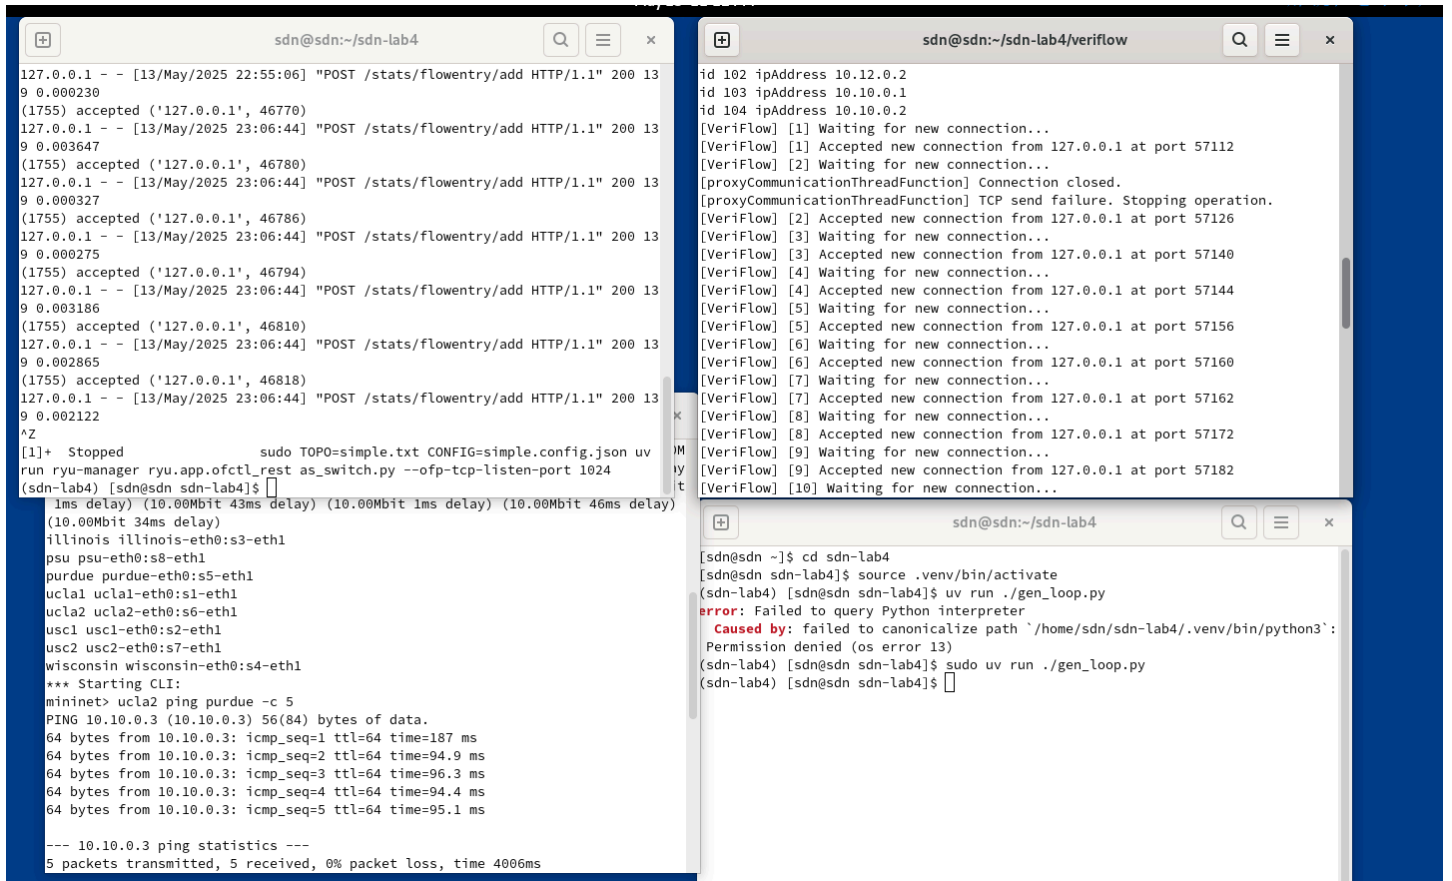
```
veriflow/VeriFlow 6633 127.0.0.1 1024 simple.txt veriflow.log
```

### 4. 启动拓扑

```
sudo ./simple.py
```

### 5. 在拓扑中ucla2 ping purdue建立连接

```
mininet> ucla2 ping purdue
```



## 6. 下发从illinois途经wisconsin到达ucla2的路径，在log文件中观察VeriFlow检测到的环路信息。

```
uv run ./gen_loop.py
```

这里发现没有相关输出，需要完成热身内容。

- 1.输出每次影响EC的数量。
- 2.打印出环路路径的信息。
- 3.进一步打印出环路对应的EC的相关信息。仅展示源MAC、目的MAC和TCP/IP五元组(即源IP、目的IP、协议、源端口、目的端口)。

### EC数目的打印

对VeriFlow.cpp的verifyRule函数进行修改

```

if(ecCount == 0)
{
    fprintf(stderr, "[VeriFlow::verifyRule] Error in rule: %s\n", rule.toString().c_str());
    fprintf(stderr, "[VeriFlow::verifyRule] Error: (ecCount = vFinalPacketClasses.size() = 0). Terminating process\n");
    exit(1);
}
else
{
    // fprintf(stdout, "\n");
    // fprintf(stdout, "[VeriFlow::verifyRule] ecCount: %lu\n", ecCount);
}

```

修改为

```
if(ecCount == 0)
{
    fprintf(stderr, "[VeriFlow::verifyRule] Error in rule: %s\n", rule.toString().c_str());
    fprintf(stderr, "[VeriFlow::verifyRule] Error: (ecCount = vFinalPacketClasses.size() = 0). Terminating process\n");
    exit(1);
}
else
{
    fprintf(stdout, "\n");
    fprintf(stdout, "[VeriFlow::verifyRule] ecCount: %lu\n", ecCount);
    fprintf(fp, "[VeriFlow::verifyRule] ecCount: %lu\n", ecCount);
}
```

### 环路路径的信息的打印

VeriFlow::traverseForwardingGraph()遍历某个特定 EC 的转发图，验证是否存在环路或黑洞。函数中的变量visited记录遍历的节点，但visited是unordered\_set类型，无法通过visited获得环路路径，因此我们需要加入新变量path用来记录环路，并在之前visited更新的位置同步更新，之后在检测到loop的时候遍历path来输出环路路径即可。

对VeriFlow.cpp的traverseForwardingGraph函数中检测环路的部分进行修改

```

...
    if(visited.find(currentLocation) != visited.end())
    {
        // Found a loop.
        fprintf(fp, "\n");
        fprintf(fp, "[VeriFlow::traverseForwardingGraph] Found a LOOP for the following packet class at\n");
        fprintf(fp, "[VeriFlow::traverseForwardingGraph] PacketClass: %s\n", packetClass.toString().c_str());

        fprintf(fp, "[VeriFlow::traverseForwardingGraph] Loop path: ");
        for(unsigned int i = 0; i < path.size(); i++)
            fprintf(fp, "%s -> ", path[i].c_str());
        fprintf(fp, "%s\n", currentLocation.c_str());

        fprintf(stdout, "\n");
        fprintf(stdout, "[VeriFlow::traverseForwardingGraph] Found a LOOP for the following packet class at\n");
        fprintf(stdout, "[VeriFlow::traverseForwardingGraph] PacketClass: %s\n", packetClass.toString().c_str());

        fprintf(stdout, "[VeriFlow::traverseForwardingGraph] Loop path: ");
        for(unsigned int i = 0; i < path.size(); i++)
            fprintf(stdout, "%s -> ", path[i].c_str());
        fprintf(stdout, "%s\n", currentLocation.c_str());

        vector<EquivalenceClass> newFaults;
        for(unsigned int i = 0; i < faults.size(); i++) {
            EquivalenceClass& fault = faults[i];
            if (fault.intersects(packetClass)) {
                vector<EquivalenceClass> diff = fault.subtract(packetClass);
                for (auto&ec : diff) {
                    newFaults.push_back(ec);
                }
            }
            else {
                newFaults.push_back(fault);
            }
        }
        faults = newFaults;
        faults.push_back(packetClass);

        return false;
    }

    visited.insert(currentLocation);
    path.push_back(currentLocation);
...

```

## 与环路对应的EC的相关信息打印

EC的基本信息打印在VeriFlow::traverseForwardingGraph()中通过调用EC类的toString()函数实现。现在将EC类的toString()函数注释掉，重新写一个toString函数即可。

```

string EquivalenceClass::toString() const
{
    char buffer[1024];

    sprintf(buffer, "[EquivalenceClass] dl_src (%s, %s), dl_dst (%s, %s)",
        ::getMacValueAsString(this->lowerBound[DL_SRC]).c_str(),
        ::getMacValueAsString(this->upperBound[DL_SRC]).c_str(),
        ::getMacValueAsString(this->lowerBound[DL_DST]).c_str(),
        ::getMacValueAsString(this->upperBound[DL_DST]).c_str());
    string retVal = buffer;
    retVal += ", ";

    sprintf(buffer, "[EquivalenceClass] nw_src (%s, %s), nw_dst (%s, %s)",
        ::getIpValueAsString(this->lowerBound[NW_SRC]).c_str(),
        ::getIpValueAsString(this->upperBound[NW_SRC]).c_str(),
        ::getIpValueAsString(this->lowerBound[NW_DST]).c_str(),
        ::getIpValueAsString(this->upperBound[NW_DST]).c_str());
    retVal += buffer;
    retVal += ", ";

    sprintf(buffer, "nw_proto(%lu, %lu)", this->lowerBound[NW_PROTO], this->upperBound[NW_PROTO]);
    retVal += buffer;
    retVal += ", ";

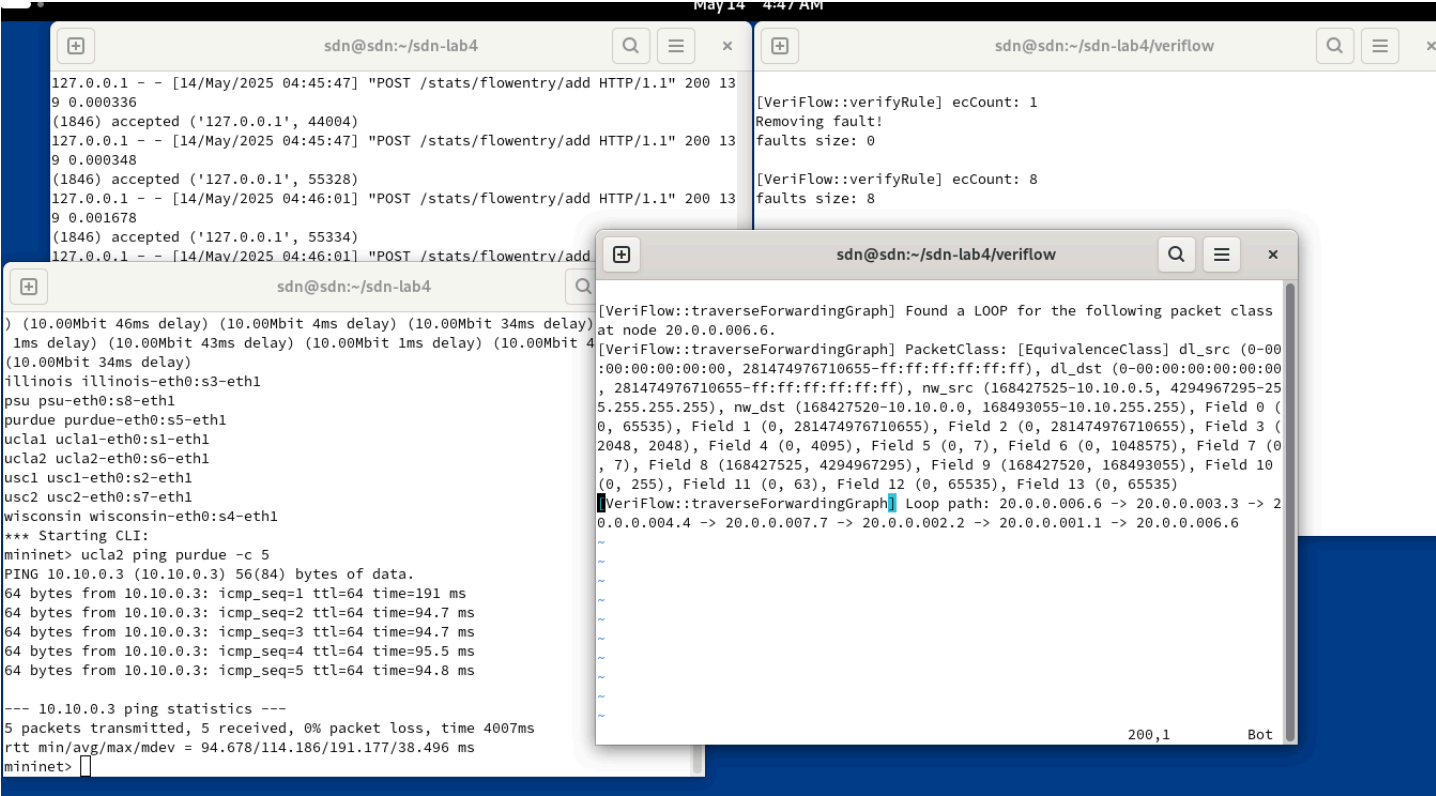
    sprintf(buffer, "tp_src(%lu, %lu)", this->lowerBound[TP_SRC], this->upperBound[TP_SRC]);
    retVal += buffer;
    retVal += ", ";

    sprintf(buffer, "tp_dst(%lu, %lu)", this->lowerBound[TP_DST], this->upperBound[TP_DST]);
    retVal += buffer;
    retVal += ", ";

    return retVal;
}

```

# 热身实验结果展示



## 四、必做任务描述及方案设计与结果

解决了刚才的转发环路问题后，你想尝试跨AS的数据包转发是否存在一些问题，于是决定测试ucla1和illinois的连通性。然而，你却发现，VeriFlow验证工具出现了一些问题。



## 1. 在自定义端口开启远程控制器，运行AS控制器程序

```
TOPO=simple.txt CONFIG=simple.config.json uv run ryu-manager  
ryu.app.ofctl_rest as_switch.py --ofp-tcp-listen-port 1024
```

## 2. 运行VeriFlow的proxy模式：

```
veriflow/VeriFlow 6633 127.0.0.1 1024 simple.txt veriflow.log
```

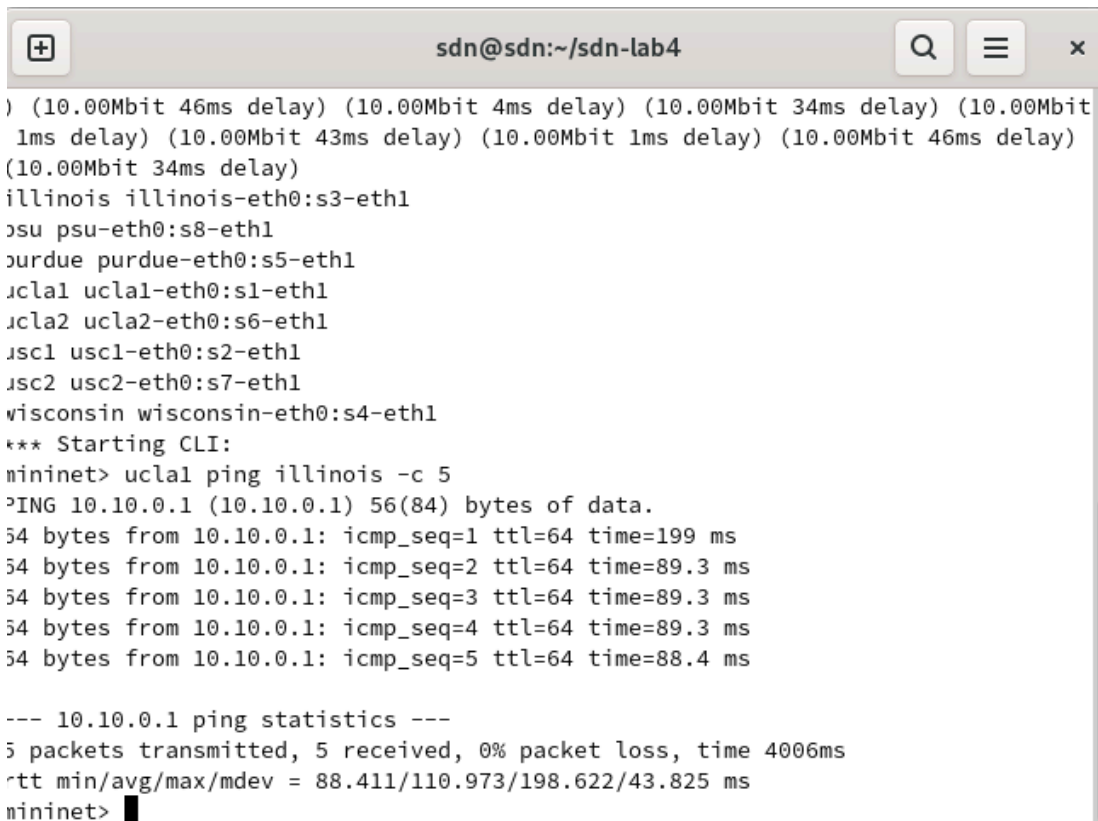
## 3. 启动拓扑

```
sudo ./simple.py
```

## 4. 建立转发路径

```
mininet> ucla1 ping illinois
```


此时，两者可以正常ping通，但veriflow提示出现了黑洞。

A terminal window titled 'sdn@sdn:~/sdn-lab4' showing the output of a ping command. The output lists network interfaces and their delays, followed by a CLI prompt. The user enters 'mininet> ucla1 ping illinois -c 5'. The terminal shows five successful ping responses from 10.10.0.1 to 10.10.0.1 with varying times. Finally, it displays ping statistics: 5 packets transmitted, 5 received, 0% packet loss, and a total time of 4006ms. The round-trip times (rtt) are listed as min/avg/max/mdev = 88.411/110.973/198.622/43.825 ms.

```
+ sdn@sdn:~/sdn-lab4
) (10.00Mbit 46ms delay) (10.00Mbit 4ms delay) (10.00Mbit 34ms delay) (10.00Mbit
1ms delay) (10.00Mbit 43ms delay) (10.00Mbit 1ms delay) (10.00Mbit 46ms delay)
(10.00Mbit 34ms delay)
illinois illinois-eth0:s3-eth1
osu psu-eth0:s8-eth1
iurdue purdue-eth0:s5-eth1
jcla1 ucla1-eth0:s1-eth1
jcla2 ucla2-eth0:s6-eth1
jsc1 usc1-eth0:s2-eth1
jsc2 usc2-eth0:s7-eth1
wisconsin wisconsin-eth0:s4-eth1
*** Starting CLI:
mininet> ucla1 ping illinois -c 5
PING 10.10.0.1 (10.10.0.1) 56(84) bytes of data.
54 bytes from 10.10.0.1: icmp_seq=1 ttl=64 time=199 ms
54 bytes from 10.10.0.1: icmp_seq=2 ttl=64 time=89.3 ms
54 bytes from 10.10.0.1: icmp_seq=3 ttl=64 time=89.3 ms
54 bytes from 10.10.0.1: icmp_seq=4 ttl=64 time=89.3 ms
54 bytes from 10.10.0.1: icmp_seq=5 ttl=64 time=88.4 ms

--- 10.10.0.1 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4006ms
rtt min/avg/max/mdev = 88.411/110.973/198.622/43.825 ms
mininet> █
```

连通性测试后的流表示意图如下：



## 分析黑洞的原因

ucla1交换机 --- (10.10.0.1/32 -> 10.12.0.1/32) ----> ucla1连接的主机

ucla1交换机 --- (10.12.0.0/16 -> 10.10.0.0/17) ----> ucla2 交换机

ucla2交换机 --- (10.12.0.1/32 -> 10.10.0.1/32) ----> illinois 交换机

illinois交换机 --- (10.12.0.1/32 -> 10.10.0.1/32) ----> illinois连接的主机

illinois交换机 --- (10.12.0.0/16 -> 10.10.0.0/17) ----> ucla2 交换机

ucla2交换机 --- (10.12.0.0/16 -> 10.10.0.0/17) ----> ucla1 交换机

主机名	交换机	dpid	IP地址	所属网段
ucla1	s1	1	10.12.0.1	10.12.0.0/16
usc1	s2	2	10.12.0.2	10.12.0.0/16
illinois	s3	3	10.10.0.1	10.10.0.0/17
wisconsin	s4	4	10.10.0.2	10.10.0.0/17
purdue	s5	5	10.10.0.3	10.10.0.0/17
ucla2	s6	6	10.10.0.4	10.10.0.0/17
usc2	s7	7	10.10.0.5	10.10.0.0/17
psu	s8	8	10.10.128.1	10.10.128.0/17

找出所有EC

### EC1

(10.10.0.1/32 → 10.12.0.1/32)

## EC2

(10.12.0.0/16 → 10.10.0.0/17)

## EC3

(10.12.0.1/32 → 10.10.0.1/32)

## EC4

(10.10.0.0/17 → 10.12.0.0/16)

对于可能的转发路径上的每个交换机，对每个EC进行分析

### ucla1交换机

(10.10.0.1/32 → 10.12.0.1/32) → ucla1主机

(10.12.0.0/16 → 10.10.0.0/17) → ucla2交换机

显然EC1 和 EC2在ucla1 交换机上都有流表项匹配

EC3 匹配第二条

EC4 在ucla1无匹配 黑洞(10.10.0.0/17 → 10.12.0.0/16)

### ucla2交换机

(10.12.0.1/32 → 10.10.0.1/32) → illinois交换机

(10.10.0.0/17 → 10.12.0.0/16) → ucla1交换机

EC3,4有匹配，EC1有匹配

EC2无匹配(10.12.0.0/16 → 10.10.0.0/17)

### illinois交换机

(10.12.0.1/32 → 10.10.0.1/32) → illinois主机

(10.10.0.0/17 → 10.12.0.0/16) → ucla2交换机

EC3,4有匹配，EC1有匹配

EC2无匹配(10.12.0.0/16 → 10.10.0.0/17)

## 5.运行uv run ./fix\_path.py

```
#!/usr/bin/env python3
from utils.flowmod import send_flow_mod

send_flow_mod(6, None, None, None, '10.10.0.0/16', None, 3)
send_flow_mod(3, None, None, None, '10.10.0.0/16', None, 1)
send_flow_mod(1, None, None, None, '10.12.0.0/16', None, 1)
```

因此给ucla2和illinois交换机下放匹配dst为10.10.0.0/16

给ucla1交换机下放匹配dst为10.12.0.0/16

即可所有EC无黑洞

新增的2个EC dst为10.10.0.0/16与dst为10.12.0.0/16

dst为10.12.0.0/16在ucla2与illinois都有匹配流表项，

而dst为10.10.0.0/16

其中dst为10.10.0.0/16的EC在ucla1交换机上并无匹配流表项。

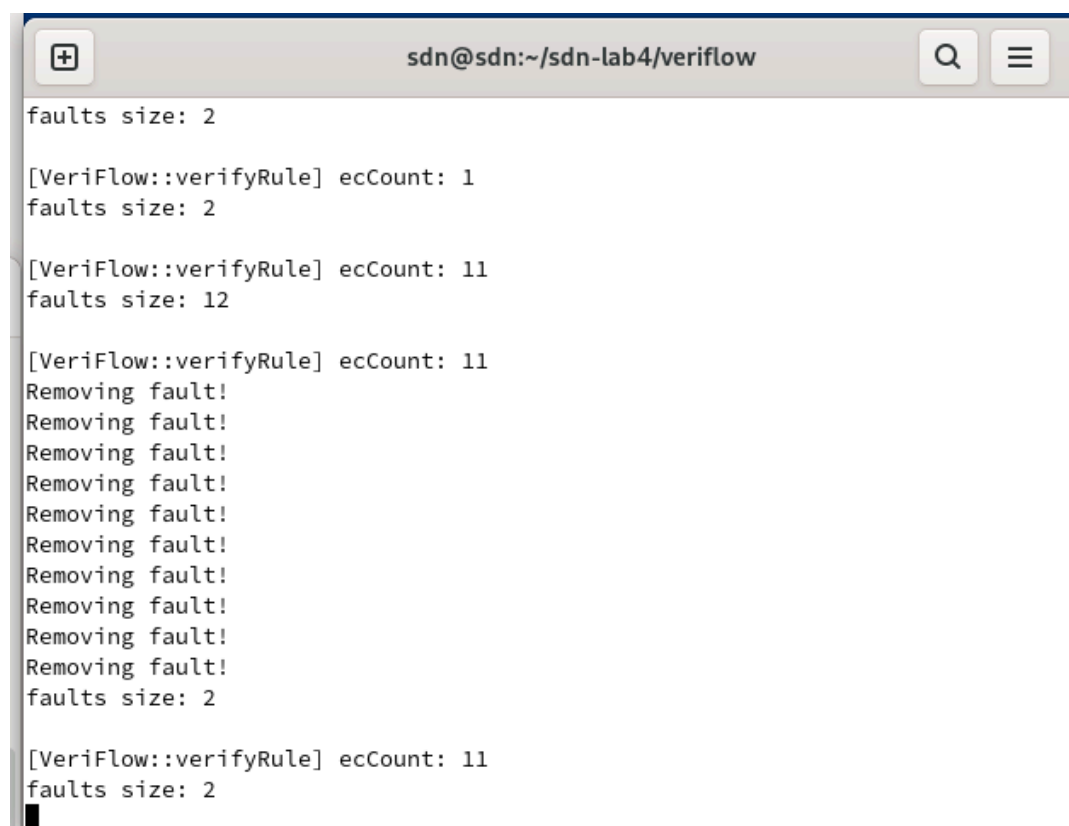
其中ucla2的dst为10.10.0.0/16的转发到端口3 s3 ill

illinois的dst为10.10.0.0/16的转发到端口1 ill

ucla1的dst为10.12.0.0/16的转发到端口1 ucla1

所以新EC也不会有黑洞。

## 6.发现veriflow依然有网络错误



```
sdn@sdn:~/sdn-lab4/veriflow
faults size: 2

[VeriFlow::verifyRule] ecCount: 1
faults size: 2

[VeriFlow::verifyRule] ecCount: 11
faults size: 12

[VeriFlow::verifyRule] ecCount: 11
Removing fault!
Removing fault!
Removing fault!
Removing fault!
Removing fault!
Removing fault!
Removing fault!
Removing fault!
Removing fault!
Removing fault!
Removing fault!
faults size: 2

[VeriFlow::verifyRule] ecCount: 11
faults size: 2
```

这是因为，在veriflow代码中，对新规则影响错误数量计算有错误。仅仅考虑了新规则的等价类完全包含错误的等价类的情况；此时，如果一个新规则的等价类覆盖了多个之前的错误，那么由于原来错误均不完全被该等价类所包含，旧的错误会依然存在；导致错误数量的计算可能有误。

## 必做任务方案设计

通过fault和packetClass是否有交集，判断是否影响到旧fault。若有交集，那fault要将除交集以外的部分删除，若删除后是空的，则将fault去除。

为此，我们要在EquivalenceClass类中实现两个方法：一是判断交集，二是求差。

### 求交集

在EquivalenceClass.h增加声明

```
bool intersects(const EquivalenceClass& other) const;
```

在EquivalenceClass.cpp中实现

```

bool EquivalenceClass::intersects(const EquivalenceClass& other) const {
    for (int i=0; i< ALL_FIELD_INDEX_END_MARKER; i++){
        if (this->upperBound[i] < other.lowerBound[i] || this->lowerBound[i] > other.upperBound[i]){
            return false;
        }
    }
    return true;
}

```

代码含义：若有一个属性无交集，则fault与EC无交集。例如src 0-1 与 src 2-3的fault和EC，无论其他属性是否有交集，都不会相交。

## 求差

在EquivalenceClass.h增加声明

```

vector<EquivalenceClass> subtract(const EquivalenceClass& other) const;

```

在EquivalenceClass.cpp中实现

```

vector<EquivalenceClass> EquivalenceClass::subtract(const EquivalenceClass& other) const {
    vector<EquivalenceClass> result;

    // 如果不相交，直接返回当前等价类
    if (!this->intersects(other)) {
        result.push_back(*this);
        return result;
    }

    // 创建一个临时等价类用于存储中间结果
    EquivalenceClass temp = *this;

    // 对每个字段进行处理
    for (int i = 0; i < ALL_FIELD_INDEX_END_MARKER; i++) {
        // 如果当前字段有重叠
        if (temp.lowerBound[i] <= other.upperBound[i] && temp.upperBound[i] >= other.lowerBound[i]) {
            // 处理下界部分
            if (temp.lowerBound[i] < other.lowerBound[i]) {
                EquivalenceClass left = temp;
                left.upperBound[i] = other.lowerBound[i] - 1;
                result.push_back(left);
            }

            // 处理上界部分
            if (temp.upperBound[i] > other.upperBound[i]) {
                EquivalenceClass right = temp;
                right.lowerBound[i] = other.upperBound[i] + 1;
                result.push_back(right);
            }

            // 更新临时等价类的边界为重叠部分
            temp.lowerBound[i] = max(temp.lowerBound[i], other.lowerBound[i]);
            temp.upperBound[i] = min(temp.upperBound[i], other.upperBound[i]);
        }
    }

    return result;
}

```

代码思路：首先调用 `this->intersects(other)` 判断两个等价类是否有交集。如果没有交集，直接返回当前等价类本身（即 `*this`），因为没有需要“减去”的部分。

创建一个临时等价类 `temp`，初始为当前等价类（`*this`），用于后续操作。

遍历每个字段，对每一个字段（如 IP、端口等，字段数量由 `ALL_FIELD_INDEX_END_MARKER` 决定）逐一处理。

对于每个字段，如果当前字段与 `other` 的对应字段有重叠（即区间有交集）：

如果当前字段的下界小于 `other` 的下界，说明有一部分在 `other` 区间左侧，这部分需要保留。

创建一个新的等价类 `left`，其上界设为 `other.lowerBound[i] - 1`，表示左侧不重叠的部分，加入结果集。

如果当前字段的下界大于 `other` 的下界，说明有一部分在 `other` 区间右侧，这部分也需要保留。

创建一个新的等价类 `right`，其下界设为 `other.upperBound[i] + 1`，表示右侧不重叠的部分，加入结果集。

最后，将 `temp` 的当前字段区间收缩为与 `other` 的重叠部分，继续处理下一个字段。

最终返回所有不与 `other` 重叠的等价类集合。

## 修改Veriflow.cpp有关部分

替换掉原来的错误代码。

原有代码（只考虑完全包含）：

```
for (unsigned int i = 0; i < faults.size(); i++) {
    if (packetClass.subsumes(faults[i])) {
        faults.erase(faults.begin() + i);
        i--;
    }
}
faults.push_back(packetClass);
```

替换为如下代码：

```
vector<EquivalenceClass> newFaults;

for (unsigned int i = 0; i < faults.size(); i++) {
    EquivalenceClass& fault = faults[i];
    if (fault.intersects(packetClass)) {
        // 1. 计算 fault - packetClass 的差集，可能得到多个新的EC
        vector<EquivalenceClass> diff = fault.subtract(packetClass);
        // 2. 把差集部分加入新的faults
        for (auto& ec : diff) {
            newFaults.push_back(ec);
        }
        // 3. 原fault被完全覆盖则不保留，部分覆盖则只保留未被覆盖部分
    } else {
        // 没有交集，原fault保留
        newFaults.push_back(fault);
    }
}
// 用新faults替换原faults
faults = newFaults;

// 最后把新检测到的错误EC加入faults
faults.push_back(packetClass);
```

这是检测到错误EC部分的替换，如果没检测到，而是其他的添加修改之类的更新，则不需向faults中push\_back，删掉最后一行即可。

代码思路就是上述提到的。

## 必做任务结果

ucla1 ping illinois 之前（运行 illinois ping wisconsin）：

```
(2076) accepted ('127.0.0.1', 49936)
127.0.0.1 - - [15/May/2025 21:43:29] "POST /stats/flowentry/add HTTP/1.1" 200 13
9 0.000411
Packet to 1
path: 1 -> 10.12.0.1
1 -> 10.12.0.1
```

```
sdn@sdn:~/sdn-lab4
(207) illinois illinois-eth0:s3-eth1
127. psu psu-eth0:s8-eth1
9 0. purdue purdue-eth0:s5-eth1
(207) ucla1 ucla1-eth0:s1-eth1
127. ucla2 ucla2-eth0:s6-eth1
9 0. usc1 usc1-eth0:s2-eth1
(207) usc2 usc2-eth0:s7-eth1
127. wisconsin wisconsin-eth0:s4-eth1
9 0. *** Starting CLI:
(207) mininet> illinois ping wisconsin -c 1
127. PING 10.10.0.2 (10.10.0.2) 56(84) bytes of data.
9 0. 64 bytes from 10.10.0.2: icmp_seq=1 ttl=64 time=106 ms

--- 10.10.0.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 105.741/105.741/105.741/0.000 ms
mininet> ucla1 ping illinois -c 1
PING 10.10.0.1 (10.10.0.1) 56(84) bytes of data.
64 bytes from 10.10.0.1: icmp_seq=1 ttl=64 time=190 ms

--- 10.10.0.1 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 189.835/189.835/189.835/0.000 ms
mininet>
```

```
sdn@sdn:~/sdn-lab4
[VeriFlow] [10] Waiting for new connection
[VeriFlow::verifyRule] ecCount: 1
faults size: 1
[VeriFlow::verifyRule] ecCount: 1
Removing fault!
faults size: 0
[VeriFlow::verifyRule] ecCount: 1
faults size: 1
[VeriFlow::verifyRule] ecCount: 1
Removing fault!
faults size: 0
[VeriFlow::verifyRule] ecCount: 1
faults size: 1
[VeriFlow::verifyRule] ecCount: 1
faults size: 5
[VeriFlow::verifyRule] ecCount: 1
Removing fault!
```

ucla1 ping illinois 之后，修复之前：

```
sdn@sdn:~/sdn-lab4
6 -> s6:2 -> 1

(2076) accepted ('127.0.0.1', 49936)
127.0.0.1 - - [15/May/2025 21:43:29] "POST /stats/flowentry/add HTTP/1.1" 200 13
9 0.000411
Packet to 1
path: 1 -> 10.12.0.1
1 -> 10.12.0.1

sdn@sdn:~/sdn-lab4
(207) illinois illinois-eth0:s3-eth1
127. psu psu-eth0:s8-eth1
9 0. purdue purdue-eth0:s5-eth1
(207) ucla1 ucla1-eth0:s1-eth1
127. ucla2 ucla2-eth0:s6-eth1
9 0. usc1 usc1-eth0:s2-eth1
(207) usc2 usc2-eth0:s7-eth1
127. wisconsin wisconsin-eth0:s4-eth1
9 0. *** Starting CLI:
(207) mininet> illinois ping wisconsin -c 1
127. PING 10.10.0.2 (10.10.0.2) 56(84) bytes of data.
9 0. 64 bytes from 10.10.0.2: icmp_seq=1 ttl=64 time=106 ms

--- 10.10.0.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 105.741/105.741/105.741/0.000 ms
mininet> ucla1 ping illinois -c 1
PING 10.10.0.1 (10.10.0.1) 56(84) bytes of data.
64 bytes from 10.10.0.1: icmp_seq=1 ttl=64 time=190 ms

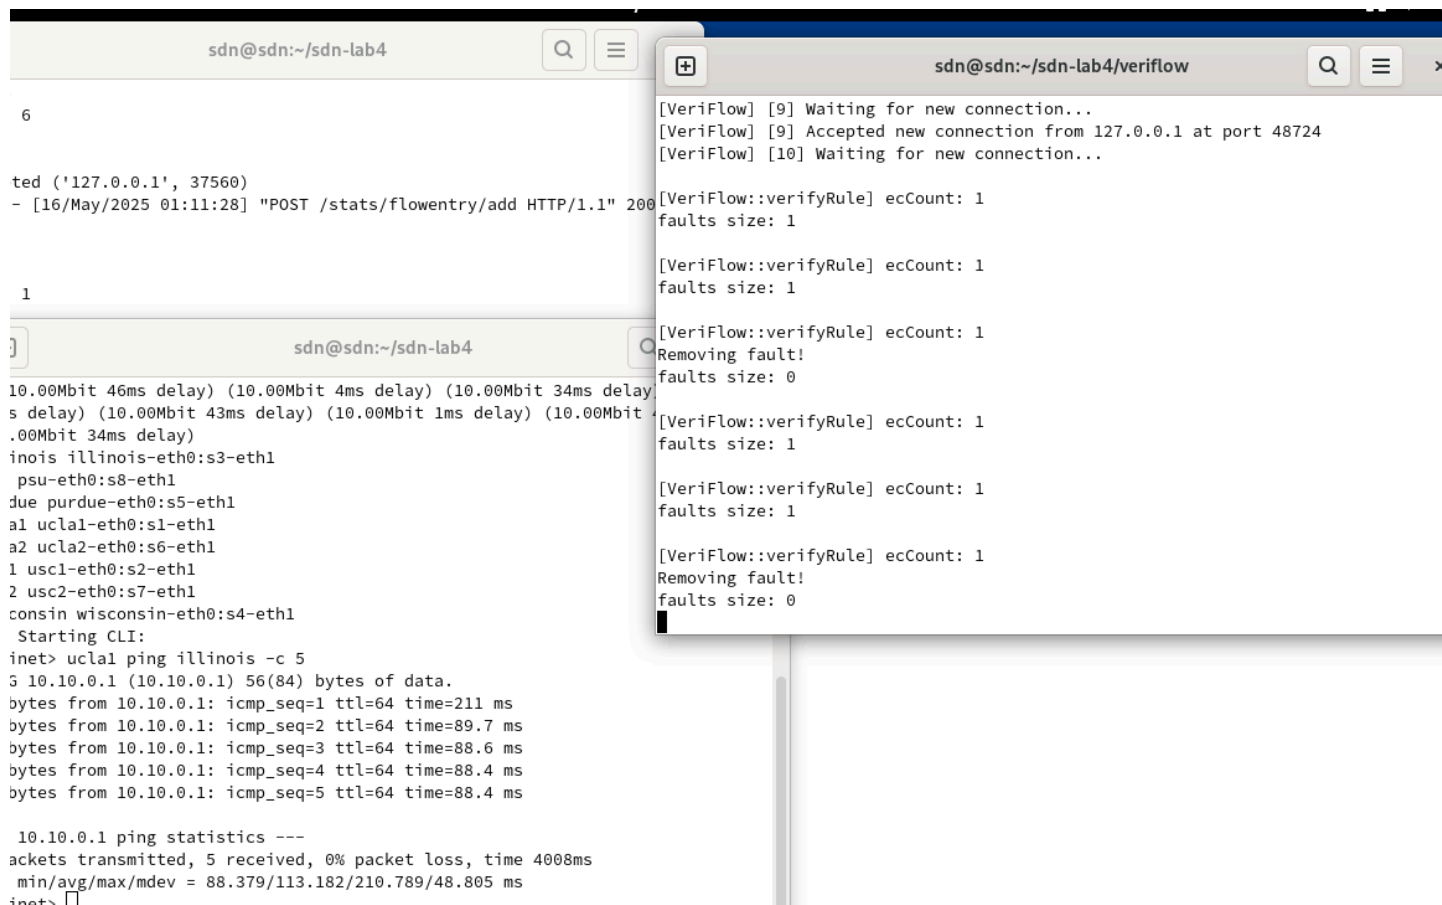
--- 10.10.0.1 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 189.835/189.835/189.835/0.000 ms
mininet>
```

```
sdn@sdn:~/sdn-lab4
[VeriFlow::verifyRule] ecCount: 1
faults size: 1
[VeriFlow::verifyRule] ecCount: 1
faults size: 5
[VeriFlow::verifyRule] ecCount: 1
Removing fault!
faults size: 4
[VeriFlow::verifyRule] ecCount: 4
faults size: 8
[VeriFlow::verifyRule] ecCount: 4
faults size: 8
[VeriFlow::verifyRule] ecCount: 1
Removing fault!
faults size: 9
[VeriFlow::verifyRule] ecCount: 16
faults size: 18
```

修复之后：







发现faults.size()为0，说明无黑洞产生了！

## 产生的黑洞问题是否会影响网络的实际使用？

不会。因为宽泛流表项匹配仅仅发生在网关交换机上，而as内的交换机到达网关交换机仍可精确匹配。这样做可以大大减少流表项数目，提高匹配效率。

## 六、实验总结

在这次lab中，我验证了 SDN 网络中流表规则配置错误（如环路、黑洞）的必然性，深刻理解动态网络环境下故障检测的必要性。通过实际案例（跨 AS 转发黑洞）发现，流表匹配范围设计与 拓扑模型一致性 对网络行为有决定性影响。通过lab，我掌握了 VeriFlow 等价类划分（EC）和 路径验证 的核心机制，理解其通过代理模式实现流表预验证的原理。掌握了VeriFlow的检测网络故障的方法，并提高了阅读工程代码、修改代码的能力。

# 附录(所有修改的代码)

In as\_switch.py

...

```
def handle_ipv4(self, msg, dpid, in_port, src_mac, dst_mac, src_ip, dst_ip, pkt_type):
```

```
....
```

```
else:
```

```
    # Handle inter-AS packet switching.
```

```
    if dpid in gateways:
```

```
        # first, if self is a gateway, send to the peer
```

```
        peer = self.routing_cfg["peers"][str(dpid)][dstnet]
```

```
        route = [dpid, peer]
```

```
        # To make matching more efficient, use source net and dst net as matching condition.
```

```
        # This might match some IPs which does not exist in the network.
```

修改:                out\_port = add\_path(route, None, None, src\_ip, dst\_ip)

```
else:
```

```
    # otherwise, send to the closest gateway
```

```
    min_delay = math.inf
```

```
    min_gw = None
```

```
....
```

```
    # To make matching more efficient, use source net and dst net as matching condition.
```

```
    # This might match some IPs which does not exist in the network.
```

修改:                out\_port = add\_path(dpid\_path, None, None, src\_ip, dst\_ip)

...

In EquivalenceClass.h

...

`class` EquivalenceClass

{

`public:`

...

加入: `bool intersects(const EquivalenceClass& other) const;`

加入: `vector<EquivalenceClass> subtract(const EquivalenceClass& other) const;`

...

};

EquivalenceClass.cpp

In EquivalenceClass.cpp

...

加入

```
bool EquivalenceClass::intersects(const EquivalenceClass& other) const {
    for (int i=0; i < ALL_FIELD_INDEX_END_MARKER; i++){
        if (this->upperBound[i] < other.lowerBound[i] || this->lowerBound[i] > other.upperBound[i]){
            return false;
        }
    }
    return true;
}

vector<EquivalenceClass> EquivalenceClass::subtract(const EquivalenceClass& other) const {
    vector<EquivalenceClass> result;

    // 如果不相交，直接返回当前等价类
    if (!this->intersects(other)) {
        result.push_back(*this);
        return result;
    }

    // 创建一个临时等价类用于存储中间结果
    EquivalenceClass temp = *this;

    // 对每个字段进行处理
    for (int i = 0; i < ALL_FIELD_INDEX_END_MARKER; i++) {
        // 如果当前字段有重叠
        if (temp.lowerBound[i] <= other.upperBound[i] && temp.upperBound[i] >= other.lowerBound[i]) {
            // 处理下界部分
            if (temp.lowerBound[i] < other.lowerBound[i]) {
                EquivalenceClass left = temp;
                left.upperBound[i] = other.lowerBound[i] - 1;
                result.push_back(left);
            }

            // 处理上界部分
            if (temp.upperBound[i] > other.upperBound[i]) {
                EquivalenceClass right = temp;
                right.lowerBound[i] = other.upperBound[i] + 1;
                result.push_back(right);
            }

            // 更新临时等价类的边界为重叠部分
            temp.lowerBound[i] = max(temp.lowerBound[i], other.lowerBound[i]);
            temp.upperBound[i] = min(temp.upperBound[i], other.upperBound[i]);
        }
    }

    return result;
}
```

```

string EquivalenceClass::toString() const
{
    char buffer[1024];

    sprintf(buffer, "[EquivalenceClass] dl_src (%s, %s), dl_dst (%s, %s)",
        ::getMacValueAsString(this->lowerBound[DL_SRC]).c_str(),
        ::getMacValueAsString(this->upperBound[DL_SRC]).c_str(),
        ::getMacValueAsString(this->lowerBound[DL_DST]).c_str(),
        ::getMacValueAsString(this->upperBound[DL_DST]).c_str());

    string retVal = buffer;
    retVal += ", ";

    sprintf(buffer, "[EquivalenceClass] nw_src (%s, %s), nw_dst (%s, %s)",
        ::getIpValueAsString(this->lowerBound[NW_SRC]).c_str(),
        ::getIpValueAsString(this->upperBound[NW_SRC]).c_str(),
        ::getIpValueAsString(this->lowerBound[NW_DST]).c_str(),
        ::getIpValueAsString(this->upperBound[NW_DST]).c_str());

    retVal += buffer;
    retVal += ", ";

    sprintf(buffer, "nw_proto(%lu, %lu)", this->lowerBound[NW_PROTO], this->upperBound[NW_PROTO]);
    retVal += buffer;
    retVal += ", ";

    sprintf(buffer, "tp_src(%lu, %lu)", this->lowerBound[TP_SRC], this->upperBound[TP_SRC]);
    retVal += buffer;
    retVal += ", ";

    sprintf(buffer, "tp_dst(%lu, %lu)", this->lowerBound[TP_DST], this->upperBound[TP_DST]);
    retVal += buffer;
    retVal += ", ";

    return retVal;
}

```

In VeriFlow.cpp

```
bool VeriFlow::verifyRule(const Rule& rule, int command, double& updateTime, double& packetClassSearchTime, double& ecCount)
{
    fprintf(fp, "[VeriFlow::verifyRule] verifying this rule: %s\n", rule.toString().c_str());

    updateTime = packetClassSearchTime = graphBuildTime = queryTime = 0;
    ecCount = 0;

    struct timeval start, end;
    double usecTime, seconds, useconds;

    gettimeofday(&start, NULL);
    // May add code in a future version to maintain a cache of forwarding graphs. This cache needs to be updated
    gettimeofday(&end, NULL);

    seconds = end.tv_sec - start.tv_sec;
    useconds = end.tv_usec - start.tv_usec;
    usecTime = (seconds * 1000000) + useconds;
    updateTime = usecTime;

    gettimeofday(&start, NULL);
    vector< EquivalenceClass > vFinalPacketClasses;
    vector< vector< Trie* > > vFinalTries;

    bool res = this->getAffectedEquivalenceClasses(rule, command, vFinalPacketClasses, vFinalTries);
    if(res == false)
    {
        return false;
    }
    gettimeofday(&end, NULL);

    seconds = end.tv_sec - start.tv_sec;
    useconds = end.tv_usec - start.tv_usec;
    usecTime = (seconds * 1000000) + useconds;
    packetClassSearchTime = usecTime;

    ecCount = vFinalPacketClasses.size();
    if(ecCount == 0)
    {
        fprintf(stderr, "[VeriFlow::verifyRule] Error in rule: %s\n", rule.toString().c_str());
        fprintf(stderr, "[VeriFlow::verifyRule] Error: (ecCount = vFinalPacketClasses.size() = 0). Terminating\n");
        exit(1);
    }
    else
    {
        fprintf(stdout, "\n");
        fprintf(stdout, "[VeriFlow::verifyRule] ecCount: %lu\n", ecCount);
        fprintf(fp, "[VeriFlow::verifyRule] ecCount: %lu\n", ecCount);
    }

    // fprintf(stdout, "[VeriFlow::verifyRule] Generating forwarding graphs...\n");
    gettimeofday(&start, NULL);
```

```

vector< ForwardingGraph* > vGraph;
for(unsigned int i = 0; i < vFinalPacketClasses.size(); i++)
{
    EquivalenceClass packetClass = vFinalPacketClasses[i];
    // fprintf(stdout, "[VeriFlow::verifyRule] [%u] ecCount: %lu, %s\n", i, ecCount, packetClass.to!
    ForwardingGraph* graph = Trie::getForwardingGraph(TP_DST, vFinalTries[i], packetClass, fp);
    vGraph.push_back(graph);
}
gettimeofday(&end, NULL);
// fprintf(stdout, "[VeriFlow::verifyRule] Generated forwarding graphs.\n");

seconds = end.tv_sec - start.tv_sec;
useconds = end.tv_usec - start.tv_usec;
usecTime = (seconds * 1000000) + useconds;
graphBuildTime = usecTime;

// fprintf(stdout, "[VeriFlow::verifyRule] Running query...\n");
gettimeofday(&start, NULL);
// Add query code here
size_t currentFailures = 0;
for(unsigned int i = 0; i < vGraph.size(); i++)
{
    unordered_set< string > visited;
    vector<string> path;
    string lastHop = network.getNextHopIpAddress(rule.location, rule.in_port);
    // fprintf(fp, "start traversing at: %s\n", rule.location.c_str());
    if(!this->traverseForwardingGraph(vFinalPacketClasses[i], vGraph[i], rule.location, lastHop, vi:
        ++currentFailures;
    }
}

fprintf(stderr, "faults size: %li\n", faults.size());
if (previousFailures > 0 && faults.size()==0) {
    fprintf(fp, "[Veriflow::verifyRule] Network Fixed!\n");
} else if (previousFailures == 0 && faults.size() > 0) {
    fprintf(fp, "[Veriflow::verifyRule] Network Broken!\n");
}
fflush(fp);
previousFailures = faults.size();
// fprintf(stdout, "[VeriFlow::verifyRule] Query complete.\n");

if(command == OFPFC_ADD)
{
    // Do nothing.
}
else if(command == OFPFC_DELETE_STRICT)
{
    Rule dummyRule = rule;
    dummyRule.type = DUMMY;

    this->removeRule(dummyRule);
}
gettimeofday(&end, NULL);

seconds = end.tv_sec - start.tv_sec;

```



```

useconds = end.tv_usec - start.tv_usec;
usecTime = (seconds * 1000000) + useconds;
queryTime = usecTime;

for(unsigned int i = 0; i < vGraph.size(); i++)
{
    delete vGraph[i];
}

return true;
}

bool VeriFlow::traverseForwardingGraph(const EquivalenceClass& packetClass, ForwardingGraph* graph, const string
{

    // fprintf(fp, "traversing at node: %s\n", currentLocation.c_str());
    if(graph == NULL)
    {
        /* fprintf(fp, "\n");
        fprintf(fp, "[VeriFlow::traverseForwardingGraph] (graph == NULL) for the following packet class
        fprintf(fp, "[VeriFlow::traverseForwardingGraph] PacketClass: %s\n", packetClass.toString().c_str()

        return true;
    }

    if(currentLocation.compare("") == 0)
    {
        return true;
    }

    if(visited.find(currentLocation) != visited.end())
    {
        // Found a loop.
        fprintf(fp, "\n");
        fprintf(fp, "[VeriFlow::traverseForwardingGraph] Found a LOOP for the following packet class at
        fprintf(fp, "[VeriFlow::traverseForwardingGraph] PacketClass: %s\n", packetClass.toString().c_str()

        fprintf(fp, "[VeriFlow::traverseForwardingGraph] Loop path: ");
        for(unsigned int i = 0; i < path.size(); i++)
            fprintf(fp, "%s -> ", path[i].c_str());
        fprintf(fp, "%s\n", currentLocation.c_str());

        fprintf(stdout, "\n");
        fprintf(stdout, "[VeriFlow::traverseForwardingGraph] Found a LOOP for the following packet class:
        fprintf(stdout, "[VeriFlow::traverseForwardingGraph] PacketClass: %s\n", packetClass.toString()

        fprintf(stdout, "[VeriFlow::traverseForwardingGraph] Loop path: ");
        for(unsigned int i = 0; i < path.size(); i++)
            fprintf(stdout, "%s -> ", path[i].c_str());
        fprintf(stdout, "%s\n", currentLocation.c_str());

        vector<EquivalenceClass> newFaults;
        for(unsigned int i = 0; i < faults.size(); i++) {
            EquivalenceClass& fault = faults[i];
            if (fault.intersects(packetClass)) {

```

```

        vector<EquivalenceClass> diff = fault.subtract(packetClass);
        for (auto&ec : diff) {
            newFaults.push_back(ec);
        }
    }
    else {
        newFaults.push_back(fault);
    }

}
faults = newFaults;
faults.push_back(packetClass);

return false;
}

visited.insert(currentLocation);
path.push_back(currentLocation);

if(graph->links.find(currentLocation) == graph->links.end())
{
    // Found a black hole.
    fprintf(fp, "\n");
    fprintf(fp, "[VeriFlow::traverseForwardingGraph] Found a BLACK HOLE for the following packet class: %s\n", packetClass.toString().c_str());
    vector<EquivalenceClass> newFaults;

    for(unsigned int i = 0; i < faults.size(); i++) {
        EquivalenceClass& fault = faults[i];
        if (fault.intersects(packetClass)) {
            vector<EquivalenceClass> diff = fault.subtract(packetClass);
            for (auto&ec : diff) {
                newFaults.push_back(ec);
            }
        }
        else {
            newFaults.push_back(fault);
        }
    }
    faults = newFaults;
    faults.push_back(packetClass);
    return false;
}

if(graph->links[currentLocation].empty() == true)
{
    // Found a black hole.
    fprintf(fp, "\n");
    fprintf(fp, "[VeriFlow::traverseForwardingGraph] Found a BLACK HOLE for the following packet class: %s\n", packetClass.toString().c_str());
    vector<EquivalenceClass> newFaults;

    for(unsigned int i = 0; i < faults.size(); i++) {

```

```

        EquivalenceClass& fault = faults[i];
        if (fault.intersects(packetClass)) {
            vector<EquivalenceClass> diff = fault.subtract(packetClass);
            for (auto&ec : diff) {
                newFaults.push_back(ec);
            }
        }
        else {
            newFaults.push_back(fault);
        }
    }

    faults = newFaults;
    faults.push_back(packetClass);

    return false;
}

graph->links[currentLocation].sort(compareForwardingLink);

const list< ForwardingLink >& linkList = graph->links[currentLocation];
list< ForwardingLink >::const_iterator itr = linkList.begin();
// input_port as a filter
if(lastHop.compare("NULL") == 0 || itr->rule.in_port == 0){
    // do nothing
}
else{
    while(itr != linkList.end()){
        string connected_hop = network.getNextHopIpAddress(currentLocation, itr->rule.in_port);
        if(connected_hop.compare(lastHop) == 0) break;
        itr++;
    }
}

if(itr == linkList.end()){
    // Found a black hole.
    fprintf(fp, "\n");
    fprintf(fp, "[VeriFlow::traverseForwardingGraph] Found a BLACK HOLE for the following packet class: %s\n", packetClass.toString().c_str());
    fprintf(fp, "[VeriFlow::traverseForwardingGraph] PacketClass: %s\n", packetClass.toString().c_str());

    vector<EquivalenceClass> newFaults;

    for(unsigned int i = 0; i < faults.size(); i++) {
        EquivalenceClass& fault = faults[i];
        if (fault.intersects(packetClass)) {
            vector<EquivalenceClass> diff = fault.subtract(packetClass);
            for (auto&ec : diff) {
                newFaults.push_back(ec);
            }
        }
        else {
            newFaults.push_back(fault);
        }
    }
}

```

```

    }
    faults = newFaults;
    faults.push_back(packetClass);

    return false;
}

if(itr->isGateway == true)
{
    // Destination reachable.
    // fprintf(fp, "[VeriFlow::traverseForwardingGraph] Destination reachable.\n");
    fprintf(fp, "\n");
    fprintf(fp, "[VeriFlow::traverseForwardingGraph] The following packet class reached destination\n");
    fprintf(fp, "[VeriFlow::traverseForwardingGraph] PacketClass: %s\n", packetClass.toString().c_str());
    vector<EquivalenceClass> newFaults;
    for(unsigned int i = 0; i < faults.size(); i++) {
        EquivalenceClass& fault = faults[i];
        if (fault.intersects(packetClass)) {
            fprintf(stderr, "Removing fault!\n");
            vector<EquivalenceClass> diff = fault.subtract(packetClass);
            for (auto&ec : diff) {
                newFaults.push_back(ec);
            }
        }
        else {
            newFaults.push_back(fault);
        }
    }
    faults = newFaults;

    return true;
}
else
{
    // Move to the next location.
    // fprintf(fp, "[VeriFlow::traverseForwardingGraph] Moving to node %s.\n", itr->rule.nextHop.c_str());

    if(itr->rule.nextHop.compare("") == 0)
    {
        // This rule is a packet filter. It drops packets.
        /* fprintf(fp, "\n");
        fprintf(fp, "[VeriFlow::traverseForwardingGraph] The following packet class is dropped\n");
        fprintf(fp, "[VeriFlow::traverseForwardingGraph] PacketClass: %s\n", packetClass.toString().c_str());
    }

    return this->traverseForwardingGraph(packetClass, graph, itr->rule.nextHop, currentLocation, visited);
}
}

```